

# healthcare capstone

May 29, 2023

```
[1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split, RandomizedSearchCV
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import cross_val_score, StratifiedKFold
from sklearn.metrics import classification_report, roc_auc_score, roc_curve, auc
```

```
[2]: # 0 for Glucose, BloodPressure, SkinThickness, Insulin, BMI doesn't make sense
      ↳and indicates missing value
dfd=pd.read_csv('health care diabetes.csv')
dfd
```

```
[2]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	
..	...	...	...	...	...	...	
763	10	101	76	48	180	32.9	
764	2	122	70	27	0	36.8	
765	5	121	72	23	112	26.2	
766	1	126	60	0	0	30.1	
767	1	93	70	31	0	30.4	

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1
..	...	...	...

763	0.171	63	0
764	0.340	27	0
765	0.245	30	0
766	0.349	47	1
767	0.315	23	0

[768 rows x 9 columns]

```
[3]: dfd.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Pregnancies           768 non-null   int64
1   Glucose               768 non-null   int64
2   BloodPressure         768 non-null   int64
3   SkinThickness         768 non-null   int64
4   Insulin               768 non-null   int64
5   BMI                   768 non-null   float64
6   DiabetesPedigreeFunction 768 non-null   float64
7   Age                   768 non-null   int64
8   Outcome               768 non-null   int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

```
[4]: # dont have null values but 0 doe
dfd.isnull().sum()
```


```
[4]: Pregnancies           0
      Glucose              0
      BloodPressure        0
      SkinThickness        0
      Insulin              0
      BMI                  0
      DiabetesPedigreeFunction 0
      Age                  0
      Outcome              0
      dtype: int64
```

```
[5]: # replace 0 values with NaN for selected columns
dfd[['Glucose', 'BloodPressure', 'SkinThickness', 'Insulin', 'BMI']] =
    ↳ dfd[['Glucose', 'BloodPressure', 'SkinThickness', 'Insulin', 'BMI']].
    ↳ replace(0, np.NaN)
```

```
[6]: # now null values show
dfd.isnull().sum().value_counts(normalize=True)
```

```
[6]: 0      0.444444
     11     0.111111
     227    0.111111
     374    0.111111
         5     0.111111
        35     0.111111
     dtype: float64
```

```
[7]: # exploring these variables using histograms
x, axes = plt.subplots(1, 5, figsize=(25,5))
sns.histplot(dfd.Glucose, ax=axes[0], color='teal')
sns.histplot(dfd.BloodPressure, ax=axes[1], color='red')
sns.histplot(dfd.SkinThickness, ax=axes[2], color='green')
sns.histplot(dfd.Insulin, ax=axes[3], color='orange')
sns.histplot(dfd.BMI, ax=axes[4], color='blue')
plt.show()
```




```
[8]: dfd2=dfd.dropna()
```


```
[9]: dfd2.shape
```

```
[9]: (392, 9)
```


```
[10]: # dropped the missing values
x, axes = plt.subplots(1, 5, figsize=(25,5))
sns.histplot(dfd2.Glucose, ax=axes[0], color='teal')
sns.histplot(dfd2.BloodPressure, ax=axes[1], color='red')
sns.histplot(dfd2.SkinThickness, ax=axes[2], color='green')
sns.histplot(dfd2.Insulin, ax=axes[3], color='orange')
sns.histplot(dfd2.BMI, ax=axes[4], color='blue')
plt.show()
```



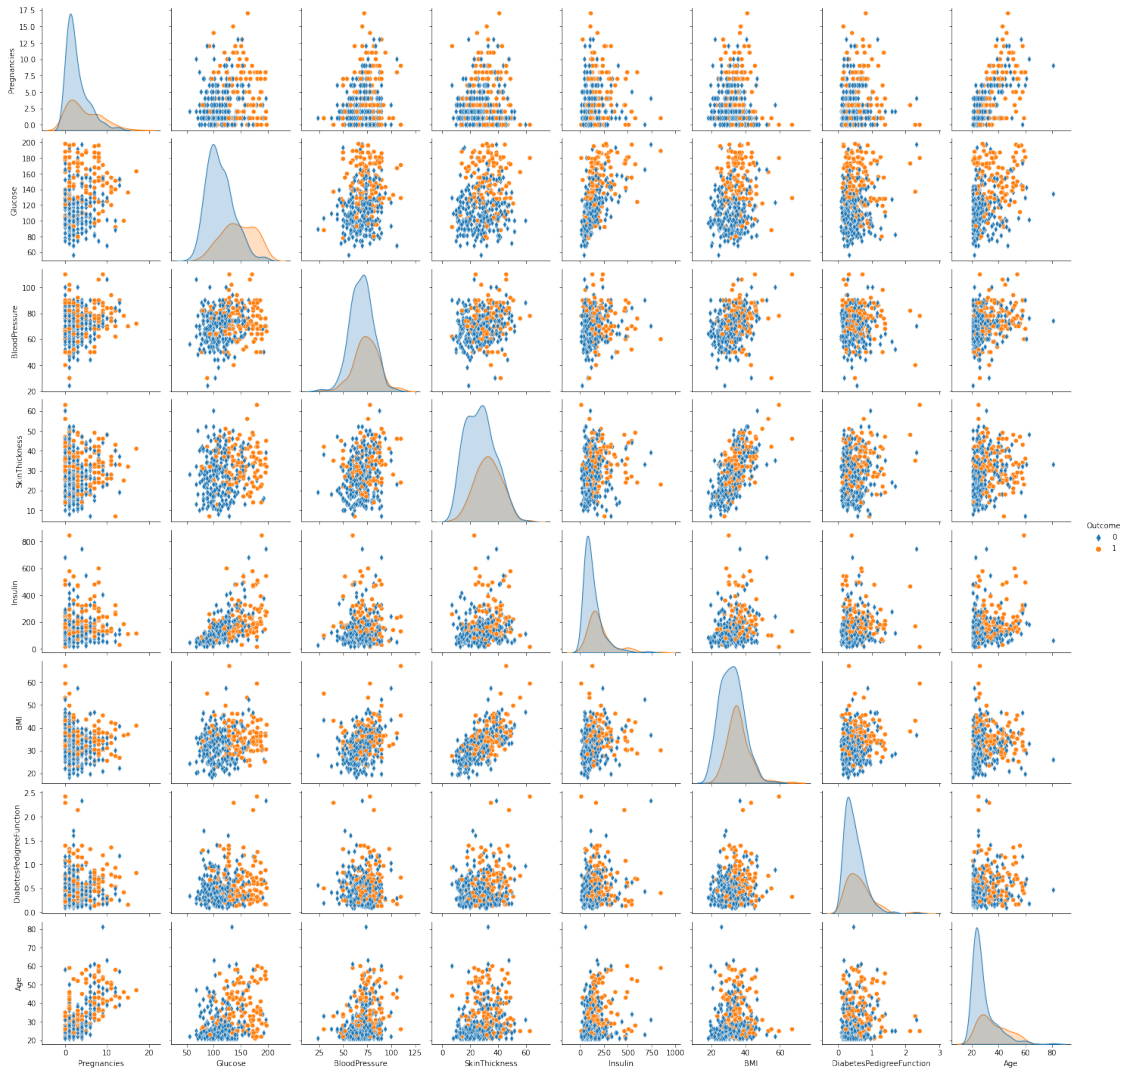
```
[11]: # count (frequency) plot describing the data types and the count of variables
sns.countplot(data=dfd, x=dfd.dtypes)
plt.title('Data Types of Variables in Diabetes Dataset')
plt.xlabel('Data Type')
plt.ylabel('Count')
plt.show()
```



```
[12]: # check the balance of the data
sns.countplot(x='Outcome', data=dfd2)
plt.title('Outcome Counts')
plt.show()
# the dataset is imbalanced, with more patients having an outcome of 0 (no
↳diabetes) than 1 (diabetes)
```



```
[13]: # scatter charts between the pair of variables to understand the relationships
sns.pairplot(dfd2, hue='Outcome', markers=['d', 'o'])
plt.show()
```



### 0.0.1 Scatterplot findings

-Age and Pregnancies shows a positive correlation. As the number of pregnancies increase, the age of the patient tends to increase as well.

-Glucose and Outcome shows a positive correlation. Patients with higher glucose levels are more likely to have diabetes.

-BloodPressure and Outcome does not show a clear correlation between the two variables.


-SkinThickness and BMI shows a positive correlation. Patients with higher BMI tend to have higher skin thickness.

-Insulin and Glucose shows a positive correlation. Patients with higher glucose levels tend to have higher insulin levels.

-BMI and Outcome shows a positive correlation. Patients with higher BMI levels are more likely to have diabetes.

-DiabetesPedigreeFunction and Age shows that patients with higher diabetes pedigree function tend to be older.

```
[14]: # correlation analysis with a heatmap
corr = dfd2.corr()
plt.figure(figsize=(12,8))
sns.heatmap(corr, annot=True, cmap='coolwarm')
plt.title('Correlation Heatmap')
plt.show()
# glucose has the biggest correlation with having diabetes
```



```
[15]: # separate the input data from the target variable, so that the model can learn
      ↳ the relationship between the two.
X = dfd2.drop('Outcome', axis=1)
# prepare data to predict outcome by setting it equal to y
y = dfd2['Outcome']
```

```
[16]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.20,
    ↪random_state=42)
print(X_train.shape, X_test.shape)
print(y_train.shape, y_test.shape)
```

```
(313, 8) (79, 8)
(313,) (79,)
```

```
[17]: # so we can refer to the variable name instead of having to type out the entire
    ↪classifier name each time we want to use it.
knn = KNeighborsClassifier()
lr = LogisticRegression(max_iter=1000)
rf = RandomForestClassifier()
svm = SVC()
```

```
[18]: # creating a dictionary named models that maps each algorithm name to its
    ↪corresponding model.
models = {'KNN': knn, 'Logistic Regression': lr, 'Random Forest': rf, 'Support
    ↪Vector Machine': svm}
```

```
[19]: # perform cross-validation on classification models using stratified 5-fold
    ↪cross-validation and report the mean accuracy and standard deviation of the
    ↪accuracy scores for each model
for name, model in models.items():
    kfold = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
    scores = cross_val_score(model, X, y, cv=kfold, scoring='accuracy')
    print(f'{name} Accuracy: {np.mean(scores)} +/- {np.std(scores)}')
```

```
KNN Accuracy: 0.7245050308341447 +/- 0.027375430575277974
Logistic Regression Accuracy: 0.7780266147354755 +/- 0.030093848040059545
Random Forest Accuracy: 0.7858487504057123 +/- 0.027227646289426977
Support Vector Machine Accuracy: 0.77036676403765 +/- 0.012148198212065897
```

```
[20]: # Train the model and make predictions
model.fit(X_train, y_train)
y_pred = model.predict(X_test)

# check if predict_proba is available
if hasattr(model, 'predict_proba'):
    y_prob = model.predict_proba(X_test)[:, 1] # Probabilities for positive
    ↪class
elif hasattr(model, 'decision_function'):
    y_scores = model.decision_function(X_test) # provides confidence scores
    y_prob = (y_scores - y_scores.min()) / (y_scores.max() - y_scores.min()) #
    ↪Normalize to [0, 1] to make them comparable to probabilities
else:
```



```

    raise AttributeError("Model does not have 'predict_proba' or
↳ 'decision_function' attribute.") # no predict_proba/decision_function
↳ attribute, no probabilities or confidence scores so raise error

```

```

[21]: # generates a classification report by comparing the predicted labels y_pred
↳ with the true labels y_test.
report = classification_report(y_test, y_pred)
# AUC score provides a summary of the model's performance to distinguish
↳ between positive and negative samples across different classification
↳ thresholds.
auc = roc_auc_score(y_test, y_prob)

```

```

[22]: # classification report is a summary of the model's performance for each class
print("Classification Report:")
print(report)
print("AUC: {:.3f}".format(auc))

```

Classification Report:


	precision	recall	f1-score	support
0	0.75	0.92	0.83	52
1	0.73	0.41	0.52	27
accuracy			0.75	79
macro avg	0.74	0.67	0.68	79
weighted avg	0.74	0.75	0.72	79

AUC: 0.770

```

[23]: # calculate roc curve
fpr, tpr, thresholds = roc_curve(y_test, y_prob)
plt.plot([0, 1], [0, 1], linestyle='--')
plt.plot(fpr, tpr, marker='.')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.show()

```



[ ]: